

TTC'16 Live Contest Case Study

Execution of dataflow-based model transformations

Antonio Garcia-Dominguez

Aston University, Birmingham, UK
a.garcia-dominguez@aston.ac.uk

Abstract.

1 Introduction

Model-to-model transformation frameworks and tools have evolved in the recent years to process larger models more efficiently. Traditional *batch* transformation engines took the entire source model(s) and produced the entire target model(s) exactly once. A small change would require going again through the entire model. Alternatively, an *incremental* transformation engine processes only the change itself, updating the target model(s) as needed.

Various approaches for incremental transformations are present in the literature: for instance, ReactiveATL [4] is an alternative execution engine of ATL which only computes and updates results on-demand as the model changes. VIATRA3 is a general purpose framework for reactive transformations [1], in which users write rules that trigger as the model changes. Streaming transformations have been studied by Cuadrado and Lara [2] with the Eclectic tool and its ATL-like language, and by David, Rath, and Varró [3] through complex event processing.

Most of these systems can be seen as complex transformations of rule-based systems into event-driven systems, or frameworks that allow for writing these event-driven systems. This case suggests studying a type of notation that may be more directly amenable to incremental execution: a graph of model-oriented primitives which generate and process streams of tuples. The notation is inspired on popular Extract-Transform-Load (ETL) tools such as Pentaho Data Integration¹ or Talend Data Integration². As data integration can be considered a form of model transformation, perhaps there might be lessons to learn from these languages. In addition to their potentially simpler incrementality, breaking rules into smaller primitives might increase the level of detail of the execution traces of the transformation.

¹<http://community.pentaho.com/projects/data-integration/>

²<https://www.talend.com/products/data-integration>

This case study presents a simplistic version of such a notation, with primitives subdivided into “minimal” and “extended” sets. Participants are tasked with writing an execution engine using their favorite approach (e.g. code generation or model interpretation) and tool, which may support batch and/or incremental transformations.

All the resources are included in the official Github repository³. These resources are mentioned in Section 2. The evaluation criteria for the provided solutions are described in Section 3. Contestants are invited to raise questions through GitHub issues should there be any unclear points in the description.

2 Case description

The present case study requires contestants to use their favorite technology to write an execution engine for a dataflow-oriented model-to-model transformation language. The transformation language has been designed with a simplistic syntax and a limited set of primitives, in order to be feasible for the time available during TTC.

The first part of the description is dedicated to introducing the language. Section 2.1 presents the abstract syntax of the language and describes the intended semantics for its various primitives. Section 2.2 explains the concrete syntaxes defined, with a HUTN-like simple textual notation and a boxes-and-arrows graphical notation. Section 2.3 gives a walkthrough of a simple transformation written in the language.

After this, two tasks are defined: one is to support the most basic primitives so the simple example shown in Section 2.3 can run. The other task is to support the rest of the primitives and enable a more complex transformation to be run.

2.1 Abstract syntax

The abstract syntax of the language has been implemented as an Ecore metamodel, which is available on the Github repository. Broadly speaking, it is based on a graph of primitives (shown in Figure 1) which may contain expression trees (Figure 2).

2.1.1 Primitives

As shown in Figure 1, a transformation is defined as a MODEL that contains a collection of connected ELEMENTS. These generally take a sequence of “rows” (a collection of key/value pairings) and process them in some way, forwarding the resulting tuples to the *target* ELEMENT. Some elements have more than one possible target in order to allow for more advanced behaviour (e.g. FILTER can split tuples into two streams according to a condition).

The primitives can be divided into several groups. First, there are the object handling primitives:

³<https://github.com/TransformationToolContest/ttc16-live-contest>

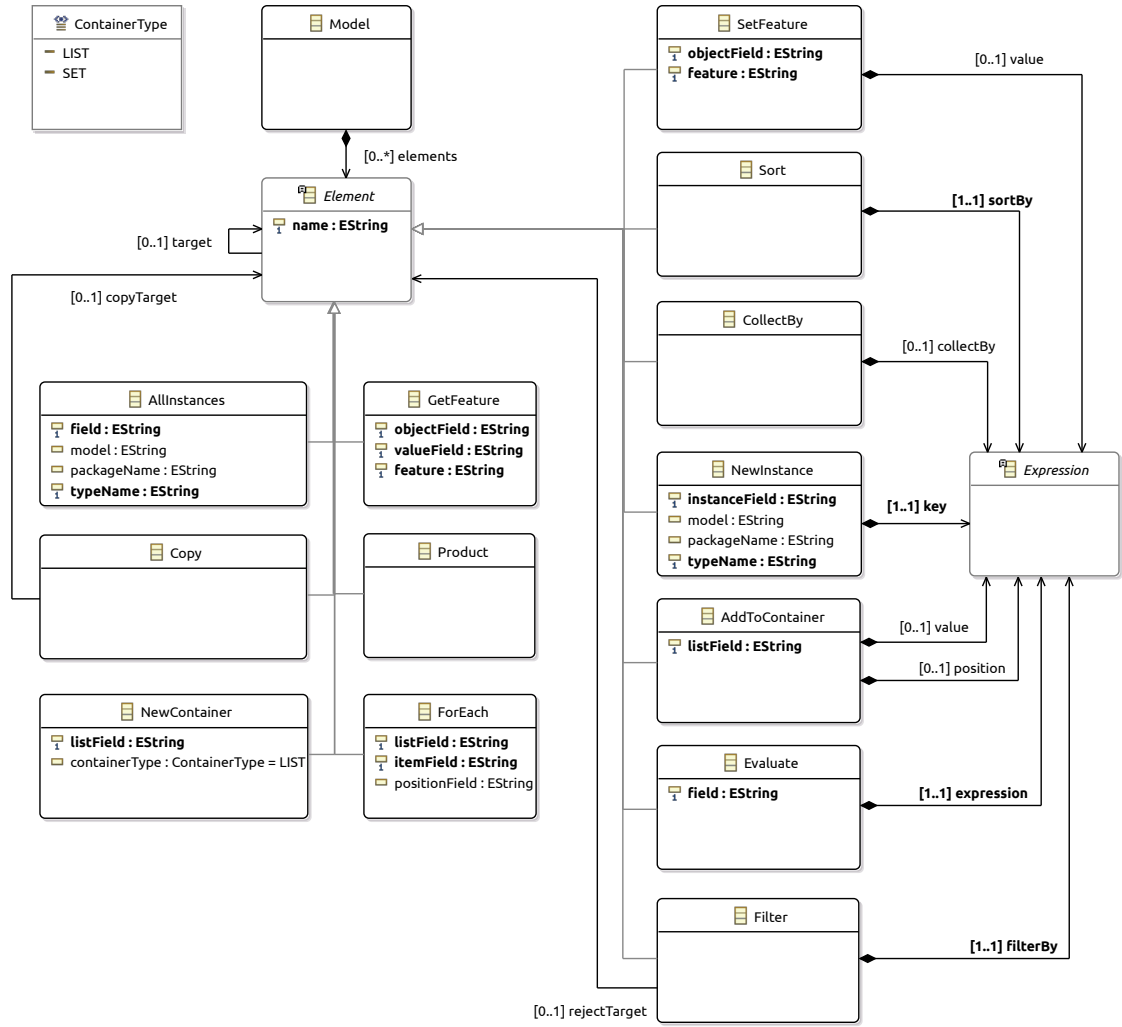


Figure 1: Abstract syntax: primitives

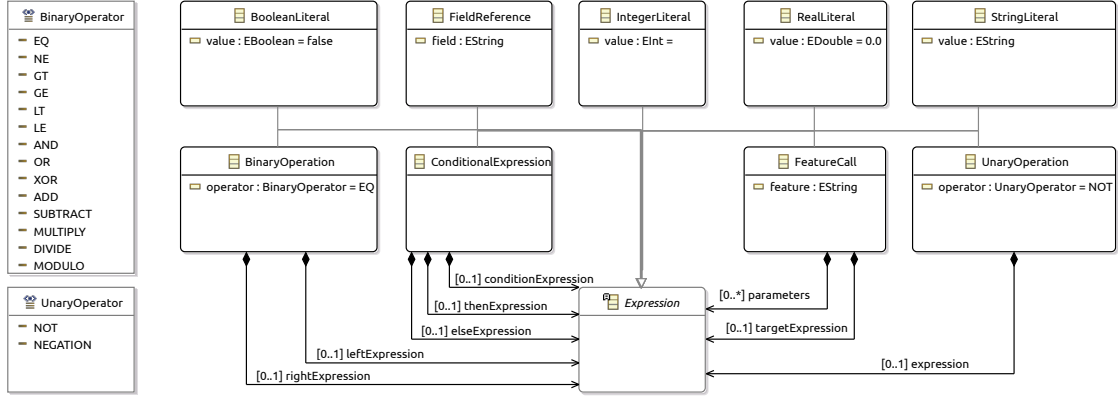


Figure 2: Abstract syntax: expressions

- **ALLINSTANCES**: for each input row r , it produces one new row r' for each instance of the specified type within the specified package and model, where the specified field will be set to that instance. Package and model are optional.

If this step has no incoming edges, it will act as it received a single empty row, producing one row per instance of the specified type from scratch. It will be very commonly the first step in most transformations.

- **NEWINSTANCE**: evaluates the *key* expression against each input tuple and tests if an instance of the same type was previously created against the same key (potentially from another **NEWINSTANCE** step). If it was not, it will set the field mentioned in *instanceField* to a new instance of the specified type and package within the specified model. Otherwise, it will set the same field to the previously created element.

The usage of a key is meant to provide a similar facility to the *equivalent()* operation in popular languages such as ATL or ETL.

- **GETFEATURE**: reads a feature (e.g. an attribute or reference) from an object in a field of the input row and places it as a field of the output row. Useful in combination with **ADDTOCONTAINER** later.
- **SETFEATURE**: sets a feature of an object to the result of a certain **EXPRESSION** on each input row.

There are primitives for computing derived values and operating on collections:

- **EVALUATE**: computes an arbitrary expression on the current input row and places it as a field on the output row.
- **NEWCONTAINER**: creates a new, empty container as a field on each input row. The container may be a set, or a list.

- **ADDTOCONTAINER**: adds an element (potentially at a certain position) to the container present in a certain field of the input row. This operation has been treated as a primitive, as it may need specific considerations for incremental processing.

More primitives are available for managing rows:

- **FILTER**: for each input row, evaluates the *filterBy* expression. Rows that produce a *true* value are sent through the *target* element, and those that produce a *false* value are sent through the *rejectTarget* element.
- **COPY**: duplicates the same input rows in the same order to the *target* and *copy-Target* elements.
- **PRODUCT**: this step is intended to receive rows from two **ELEMENTS** and generate the cartesian product of the two sets of rows (e.g. produce all pairs of rows from both elements).
- **SORT**: reads in all the rows and then sorts them according to the values produced by the *sortBy* expression.
- **FOREACH**: for each input row with a collection on a certain field, it produces as many rows as elements in that collection, setting a certain field to each of its elements.
- **COLLECTBY**: for each sequence of contiguous rows with the same value for *collectBy*, it will produce one row replacing all fields with collections of their respective values in the sequence.

While these are quite a few primitives, not all of them need to be implemented to run the transformations proposed in this case study. The required primitives will be mentioned later on.

2.1.2 Expressions

As mentioned above, some of these **ELEMENTS** can embed **EXPRESSIONS** in order to compute derived values, sorting/grouping keys or filtering conditions. Figure 2 shows the available elements for the small expression language. The objective of the language is to be side-effect free as much as possible. An expression is a tree of elements of various types:

- Literals of a certain type, e.g. **BOOLEANLITERAL** for boolean values and so forth.
- **FIELDREFERENCES** to a certain field within the row, by name.
- **UNARYOPERATIONS**, which take the result of a subexpression and apply to it logical negation (*NOT* in **UNARYOPERATOR**) or arithmetic negation (*NEGATION*).

- **BINARYOPERATIONS**, which combine the result of two subexpressions in various ways. The language supports equality comparison (*EQ* in **BINARYOPERATOR**), inequality (*NE*), greater than (*GT*), greater or equal (*GE*), less than (*LT*), less or equal (*LE*), logical AND with shortcircuit, logical OR with shortcircuit, logical XOR, numeric addition / string concatenation (*ADD*), subtraction, multiplication, division or the modulo operation.
- **FEATURECALLS**, which access a certain property or method within an object contained in the present row. If it is a method call, it will include a list of parameter subexpressions. The transformations to be run assume that “a.x” will be translated to “a.eGet(feature x)” in EMF terms, and that “a.eClass” and “a.eContainer” will also be available.
- **CONDITIONALEXPRESSIONS**, inspired by the ternary operator in C/C++ and by the *if..then..else* expression in Python. If the conditional expression produces a *true* value it will compute and return the result of the *thenExpression*, otherwise it will use the *elseExpression*.

2.2 Concrete syntax

Explain the textual (Xtext-based) and graphical (Sirius-based) notations. Doesn't need to be too detailed, as contestants won't actually write transformations.

2.3 Example: Families to Persons

Explain the Families2Persons example as described through the notations.

2.4 Task 1: base primitives

families2persons could be implemented, then tree2graph could be the “extra” transformation for soundness.

2.5 Task 2: extended primitives

flowchart2html could be implemented, then class2rdb could be the “shadow” transformation for soundness.

3 Evaluation

Correctness (w/shadow transformation with base primitives revealed during conference), performance, extensibility.

References

- [1] Gábor Bergmann, István Dávid, Ábel Hegedüs, Ákos Horváth, István Ráth, Zoltán Ujhelyi, and Dániel Varró. Viatra 3: A Reactive Model Transformation Platform. In Dimitris Kolovos and Manuel Wimmer, editors, *Theory and Practice of Model Transformations*, number 9152 in Lecture Notes in Computer Science, pages 101–110. Springer International Publishing, July 2015.
- [2] Jesús Sánchez Cuadrado and Juan de Lara. Streaming Model Transformations: Scenarios, Challenges and Initial Solutions. In Keith Duddy and Gerti Kappel, editors, *Theory and Practice of Model Transformations*, number 7909 in Lecture Notes in Computer Science, pages 1–16. Springer Berlin Heidelberg, June 2013. DOI: 10.1007/978-3-642-38883-5_1.
- [3] István Dávid, István Ráth, and Dániel Varró. Streaming Model Transformations By Complex Event Processing. In Juergen Dingel, Wolfram Schulte, Isidro Ramos, Silvia Abrahão, and Emilio Insfran, editors, *Model-Driven Engineering Languages and Systems*, number 8767 in Lecture Notes in Computer Science, pages 68–83. Springer International Publishing, September 2014.
- [4] Massimo Tisi, Salvador Martínez, Frédéric Jouault, and Jordi Cabot. Lazy execution of model-to-model transformations. In *Model Driven Engineering Languages and Systems*, pages 32–46. Springer, 2011.